

Things to know while choosing a Deep-Learning Library

Saurabh Agarwal

Machine Learning Engineer

MAD Street Den

Why DL libraries ?

- Quick Prototyping
- Reducing Engineering and Focus on building algorithms.




What can I choose from ? (there are many more)



Caffe2

Caffe


theano



Keras

Lasagne

dy/net



Chainer

PYTORCH

Confused ?

- **Library - Library everywhere not a clue which one to use.**
- **Everybody claims theirs is faster or easier to use.**



Structure of the talk

1. Differences and what those differences mean ?
2. State of some popular ML libraries
3. Example scenarios with intended use cases.

Take a step back ?

- Fundamental Questions -
 - How do they differ ?
 - What benefits do I get because of these differences ?



Language the framework is written

- Python - (Theano)
 - C++ - (Caffe)
 - Lua - (Torch)
-
1. Frameworks written in Python usually are slower. But easy to modify for a custom change.
 2. Most of them have Python-API. So you get the speed of low-level language and ease of coding in python

Community Support and Documentation

- How active is the community ?
- Is the project under active development ?
- Is the project stable ?
- Are there any big companies supporting it ?
- Documentation ?

Graph ?

- **What is a graph ?**
 - Typically the frameworks maintain a computational graph.
 - It specifies the sequence of operations to perform.
- **Types of Graph:**
 - **Static - Define and Run**
 - **Dynamic - Define by Run**

More on Static Graph

- First a graph is defined in host language (ex- Python) .
- A second interpreter that executes the graph.
- Pros:
 - Optimization
 - Utilization of SIMD (single instruction , multiple data)
 - Overall-Faster at runtime
- Cons:
 - Slow compile time
 - Difficult to debug, because of two interpreter structure
 - Almost impossible to build certain recursive neural net architecture

More on Dynamic Graphs

- Every iteration a new graph is built
- The states of older variables is saved in the graph for next iteration
- Pros:
 - Complex recursive architectures
 - Easy Debugging - standard python tools like pdb work
- Cons:
 - Slower at runtime
 - High memory consumption

Using a High-Level API

- Toolkits like Keras, Lasagne
- Provide simplified API's to allow quick prototyping
- Use other libraries in the backend .
- Pros:
 - Helper functions to handle grunt work
- Cons:
 - Another level of Abstraction.

Inference Speed on GPU

- Speed usually is almost same because of CuDNN
- The delta can come because of graph creation overhead

State of Some Popular Libraries

Tensorflow

- **The Good things:**
 - **Excellent Community support**
 - **Stable Release**
 - **Under Active Development**
 - **Google's Support**
 - **Lot pretrained Networks and Open Source implementation of Networks**
- **The not so Good things:**
 - **Static Graph Library**
 - **Overwhelming**

Caffe

- **The Good things**
 - **Very Fast**
 - **Lot of Pretrained Networks**
 - **Great Community Support**
- **The not so Good things**
 - **C++, very difficult to modify**

Theano

- **Good Things**
 - **Good Community Support**
 - **Python mostly**
 - **Allows you to do a lot of crazy stuff**
- **The not so Good things**
 - **Slow at inference**
 - **Single GPU**
 - **Graph is very low level**

Chainer

- **The Good things**
 - **First to support Dynamic Graphs**
 - **All benefits of Dynamic Graphs**
 - **Highly scalable multi-gpu version**
- **The not so Good things**
 - **Slower than other dynamic libraries**
 - **Small Community**

Dynet

- **The Good things**
 - **Almost as fast as static libraries on CPU**
 - **Python bindings**
 - **Very Pythonic API Structure**
- **The not so Good things**
 - **Small user group**
 - **Not backed by any major company**

PyTorch

- **The Good things:**
 - **Torch without Lua**
 - **Dynamic Graph**
 - **Promises Numpy on GPU**
 - **Amazing Community**
- **The not so Good things**
 - **Still in Beta**
 - **Can be slow at times**

Keras

- **Good things**
 - **Amazing Community Support**
 - **Lot of built in functionalities**
 - **Simple API**
- **Not so Good things**
 - **Depends on the library used in the backend**

MXNet

- **Good things**
 - **Apache Supported Project**
 - **Fast**
 - **Under Active Development**
 - **Great Documentation**
 - **AWS Supports**

**What works for one
might not work for
another**

Intended Usage

- **Research vs Production**
 - **For Research :**
 - Ability for rapid prototyping
 - Ease of debugging
 - Speed takes a backseat
 - **For Production**
 - Speed takes prime importance
 - Ability of Library to Scale

Example Scenario

Beginner

- Easy API
- Lot of support for quick saving and loading of models
- Good Documentation
- Community Support
- Better off using toolkits like KERAS

Researcher

- Likes to prototype quickly to test new and crazy hypothesis
- Ease of debugging
- Doesn't care much about speed
- Dynamic Graph libraries like Chainer and Pytorch make more sense

Practitioner in Industry

- Cares about speed
- Availability of pre-trained models
- Adoption by other companies.
- Static Graph libraries like Caffe, Tensorflow and Caffe2 makes more sense
- Also if you do deployments on a GPU then the differences can be negligible, CUDNN at work
- A hybrid approach is usually a good idea. Development in a dynamic graph library and deployment in static libraries

Multi Platform Deployment

- Tensorflow shines over everybody.

General Advice

- Get started with something you feel comfortable.
- Move around as and when situation demands
- Community Support
- You get hang of one type everything else makes sense

THANK YOU!!